

Lauren Fromm
404751250

Lab 5 Log

First, I need to create the tr2b.c file
`$ touch tr2b.c`

Inside this file, I first check that there is a correct amount of arguments inputted (3). Then, I make sure that the 'from' argument and the 'to' argument are the same length. Then, I need to make sure that none of the characters in the 'from' argument are repeated, so I go through every character and check that it has not already been listed. After all my checks, I can go through every character of the input file using `getchar()`, and if it matches one of the characters in the 'from' argument, I replace it with the character that it correlates to in the 'to' argument. I output the everything with `putchar()`.

Next, I create my tr2u.c file:
`$ touch tr2u.c`

I use the same logic when creating this file, except I need to use `read` and `write`. I go through every character of the input file using a while loop that reads a character every time. When I go through every character of the input file, I put each character in an array and then test it to see if it needs to be replaced. If it does need to be replaced, I use `write` to write it out to output.

To get a big file, I use:
`$ head --bytes=5000000 /dev/urandom > output.txt`
which puts a 5,000,000 byte file into output.txt

I want to use the `strace` command to find out what system call each file uses. I use the `-c` flag to show time, calls, and errors for each system call and report a summary on program exit. I first copy one file to another by outputting the final result into files called `outputtr2u` and `outputtr2b`.

My `strace` command then looks like:
`$ strace -c ./tr2b 'ab' 'yz' < output.txt > outputtr2b`

which outputs:

% time	seconds	usecs/call	calls	errors	syscall
0.00	0.000000	0	2		read
0.00	0.000000	0	1		write
0.00	0.000000	0	2		open
0.00	0.000000	0	2		close
0.00	0.000000	0	4		fstat
0.00	0.000000	0	9		mmap
0.00	0.000000	0	4		mprotect
0.00	0.000000	0	1		munmap
0.00	0.000000	0	1		brk
0.00	0.000000	0	1	1	access
0.00	0.000000	0	1		execve
0.00	0.000000	0	1		arch_prctl
100.00	0.000000		29	1	total

I then run the command with the unbuffered file:

```
$ strace -c ./tr2u 'ab' 'yz' < output.txt > outputtr2u
```

which outputs:

% time	seconds	usecs/call	calls	errors	syscall
54.96	21.183129	4	5000000		write
45.03	17.356050	3	5000002		read
0.00	0.000059	8	7		mmap
0.00	0.000027	7	4		mprotect
0.00	0.000014	7	2		open
0.00	0.000010	10	1		munmap
0.00	0.000007	4	2		close
0.00	0.000007	7	1	1	access
0.00	0.000007	7	1		execve
0.00	0.000006	3	2		fstat
0.00	0.000004	4	1		brk
0.00	0.000003	3	1		arch_prctl
100.00	38.539323		10000024	1	total

The unbuffered version that uses read and write has a lot more system calls than the buffered putchar and getchar file. Because of this, tr2u uses a lot more time.

If i want the output to be directly copied into terminal, I delete the > outputtr2b and > outputtr2u:

```
$ strace -c ./tr2b 'ab' 'yz' < output.txt
```

Which gives:

time	seconds	usecs/call	calls	errors	syscall
0.00	0.000000	0	2		read

0.00	0.000000	0	2	write
0.00	0.000000	0	2	open
0.00	0.000000	0	2	close
0.00	0.000000	0	4	fstat
0.00	0.000000	0	9	mmap
0.00	0.000000	0	4	mprotect
0.00	0.000000	0	1	munmap
0.00	0.000000	0	1	brk
0.00	0.000000	0	1	1 access
0.00	0.000000	0	1	execve
0.00	0.000000	0	1	arch_prctl

100.00	0.000000		30	1 total

Then I do the same for the tr2u:

```
$ strace -c ./tr2u 'ab' 'yz' < output.txt
```

There is a very long list of random characters outputted and then the final summary:

time	seconds	usecs/call	calls	errors	syscall
54.41	22.161857	4	5000000		write
45.59	18.569249	4	5000002		read
0.00	0.000000	0	2		open
0.00	0.000000	0	2		close
0.00	0.000000	0	2		fstat
0.00	0.000000	0	7		mmap
0.00	0.000000	0	4		mprotect
0.00	0.000000	0	1		munmap
0.00	0.000000	0	1		brk
0.00	0.000000	0	1	1	access
0.00	0.000000	0	1		execve
0.00	0.000000	0	1		arch_prctl

100.00	40.731106		10000024		1 total

Again, the unbuffered file has a significantly more amount of system calls than the buffered file. The reason for this is that the unbuffered file is making a system call everytime it has to read or write a byte, and since the file is so huge, theres an obvious discrepency between the two files.

I then use the time command to see the different times each file has.

```
$ time ./tr2b 'ab' 'yz' < output.txt
```

which outputs:

```
real 0m0.002s
user 0m0.000s
sys 0m0.001s
```

```
$time ./tr2u 'ab' 'yz' < output.txt  
which outputs a long list of random characters and then:  
real 0m5.167s  
user 0m0.317s  
sys 0m4.814s
```

The unbuffered version again takes a lot longer than the buffered version to execute.